

# 03\_Linear\_Regression\_Concise

December 23, 2023

```
[31]: import torch
```

## 1 Concise Implementation of Linear Regression

- This notebook shows how linear regression can be implemented using higher-level functionalities provided by PyTorch.

## 2 Creating the Dataset

```
[32]: # Generating the Dataset
def synthetic_data(w, b, num_examples):
    """Generate  $y = Xw + b + \text{noise}$ ."""
    X = torch.normal(0, 1, (num_examples, len(w)))
    y = torch.mm(X, w) + b
    y += torch.normal(0, 0.01, y.size())
    return X, y

true_w = torch.tensor([[2], [-3.4]])
true_b = torch.tensor(4.2)
features, labels = synthetic_data(true_w, true_b, 1000)
```

## 3 Generating batches

- We will generate batches using high-level functionalities provided by PyTorch classes.
- The abstract class `Dataset` requires implementing a function called `__len__` (called by Python's `len` function) and a function called `__getitem__`, which enables indexing, iterating, and slicing.
- The class `TensorDataset` implements the `Dataset` interface. The constructor of this class accepts an arbitrary number of tensors whose first dimension have the same length. When indexed, a `TensorDataset` returns a tuple with the corresponding elements from each of these tensors.
- The class `DataLoader` enables iterating through minibatches of a `Dataset`.

```
[33]: dataset = torch.utils.data.TensorDataset(features, labels) # Creates a
      ↪ `TensorDataset`
      print(dataset[0]) # The first example in our dataset

      print()

      batch_size = 10
      data_iter = torch.utils.data.DataLoader(dataset, batch_size, shuffle=True) #
      ↪ Creates a `DataLoader`
      next(iter(data_iter)) # Creates an iterator from the `DataLoader` and requests
      ↪ the first element
```

```
(tensor([0.3225, 1.2098]), tensor([0.7241]))
```

```
[33]: [tensor([[ 0.8063, -0.8220],
               [-0.6907,  0.0055],
               [ 0.9742,  0.0982],
               [-0.8409,  0.3336],
               [ 2.4879,  1.4182],
               [-0.1178,  0.0546],
               [-1.4182, -2.2553],
               [ 0.6778,  0.2994],
               [-0.0147,  1.4285],
               [-0.3188, -0.2341]]),
      tensor([[ 8.5917],
               [ 2.7999],
               [ 5.8137],
               [ 1.3876],
               [ 4.3666],
               [ 3.7818],
               [ 9.0093],
               [ 4.5373],
               [-0.6812],
               [ 4.3846]])]
```

## 4 Defining the Model

- Recall that a linear model can be interpreted as a very simple neural network.
- We will use the neural network functionalities provided by PyTorch to define our model, effectively creating and training a neural network!
- In neural network terms, we need to use a fully-connected linear layer, which is implemented by the `Linear` class in PyTorch.

```
[34]: num_of_inp = 2 # Number of inputs to the layer
      num_of_out = 1 # Number of outputs from the layer
      net = torch.nn.Linear(num_of_inp, num_of_out) # Creates our model (a neural_
      ↪network with a fully-connected linear layer and a single output)
```

## 5 Initializing Model Parameters

- We will initialize the parameters of the neural network as in the previous notebook.
- We can modify the parameters by accessing the weights (`net.weight.data`) and the bias (`net.bias.data`).
- Accessing the data of the tensors eliminates the need for `torch.no_grad`.
- The in-place methods `normal_` and `fill_` can be used to overwrite parameter values.

```
[35]: net.weight.data.normal_(0, 0.01); # Each weight is sampled from a normal_
      ↪distribution with mean 0 and standard deviation 0.01.
      net.bias.data.fill_(0); # The bias is initialized to 0.
```

## 6 Defining the Loss Function

- The `MSELoss` class computes the mean squared error.

```
[36]: loss = torch.nn.MSELoss()
```

## 7 Defining the Optimization Algorithm

- Stochastic gradient descent is implemented by the `SGD` class.
- The constructor of this class requires a list of parameters to be optimized (which can be obtained through `net.parameters()`) and accepts some hyperparameters (such as the learning rate `lr`).

```
[37]: optimizer = torch.optim.SGD(net.parameters(), lr=0.5)
```

## 8 Training Loop

- During each **epoch**:
  - Execute one iteration per minibatch.
  - During each iteration:
    - \* Obtain the minibatch.
    - \* Compute predictions and loss using the current model (**forward pass**).
    - \* Compute the gradients of the loss with respect to model parameters (**backward pass**).
    - \* Update the model parameters.

```
[38]: print('\nInitial parameters:')
      print(net.weight)
      print(net.bias)

      print()

      num_epochs = 3
      for epoch in range(num_epochs):
          for X, y in data_iter: # Minibatch: `X` and `y`
              y_hat = net(X) # Prediction for the minibatch
              l = loss(y_hat, y) # Loss for the minibatch
              optimizer.zero_grad() # Zeroes the gradient stored inside each parameter
              l.backward() # Computes gradient of `l` with respect to parameters
              optimizer.step() # Updates each parameter based on the gradient stored
                               ↪ inside it.

              # After each epoch, computes the loss for the entire training dataset
              l = loss(net(features), labels)
              print(f'Epoch {epoch + 1}. Loss: {l:f}')

      print('\nLearned parameters:')
      print(net.weight)
      print(net.bias)

      print('\nTrue parameters:')
      print(true_w)
      print(true_b)
```

```
Initial parameters:
Parameter containing:
tensor([[0.0114, 0.0151]], requires_grad=True)
Parameter containing:
tensor([0.], requires_grad=True)
```

```
Epoch 1. Loss: 0.000142
Epoch 2. Loss: 0.000112
Epoch 3. Loss: 0.000136
```

```
Learned parameters:
Parameter containing:
tensor([[ 1.9988, -3.3948]], requires_grad=True)
Parameter containing:
tensor([4.2021], requires_grad=True)
```

```
True parameters:
tensor([[ 2.0000],
```

```
[-3.4000]])  
tensor(4.2000)
```

## 9 Evaluation

- Because we created the dataset, we can evaluate our success by comparing the true parameters with the learned parameters.

## 10 [Storing this notebook as a pdf]

```
[39]: %%capture  
from google.colab import drive  
drive.mount('/content/gdrive', force_remount=True)  
  
!sudo apt-get install texlive-xetex texlive-fonts-recommended  
↪texlive-plain-generic  
  
# Set the path to this notebook below (add \ before spaces). The output `pdf`  
↪will be stored in the corresponding folder.  
!jupyter nbconvert --to pdf /content/gdrive/My\ Drive/Colab\ Notebooks/nnd1/  
↪week_03/lecture/03_Linear_Regression_Concise.ipynb  
  
# If having issues, save this notebook (File > Save) and restart the session  
↪(Runtime > Restart session) before running this cell. To debug, remove the  
↪first line (%%capture).
```